



INVESTIGATION REPORT



Establishing a Workflow for Real-time Global Illumination and Vector Graphics in Enhancing Dynamic Game Environments

By

Cheng Yi Heng

TP058994

APU3F2408CGD

A report submitted in partial fulfilment of the requirements for the degree of
BSc (Hons) Computer Games Development
at Asia Pacific University of Technology and Innovation.

Supervised by Mr. Jacob Sow Tian You

2nd Marker: Dr. Tan Chin Ike

27th November 2024

Acknowledgement

Abstract

Achieving visually rich and interactive content in real-time without compromising performance is a key aspect of immersive gameplay. This project seeks to establish a streamlined workflow that incorporates real-time global illumination (GI) and compute-centric vector graphics (VG) into dynamic game environments in a way that is accessible and adaptable for game developers. These integrations will help improve visual appeal and interactive feedback in gameplay. Our method utilizes radiance cascades for enabling real-time GI, Vello for rendering dynamic VG in real-time, and Typst for VG content creation. Using purposive sampling, this study targets developers in the game development industry through the use of questionnaires. It is aimed to gather insights on the importance of GI and VG as well as the level of integrations in state-of-the-art game engines. This paper provides an in-depth look at the challenges and potential methods for integrating these technologies into game development, with an emphasis on their impact on interactive and adaptable content creation. Our approach contributes to advancing infrastructure and fostering innovation, aligning with the goals of Sustainable Development Goal (SDG) 9.

Keywords — radiance cascades, indirect lighting, typesetting, markdown, interactivity, dynamic content

Contents

Acknowledgement	2
Abstract	3
CHAPTER 1: INTRODUCTION	7
1.1. Problem Background	8
1.1.1. Limitations of scalability in current real-time global illumination techniques	8
1.1.2. Advantages of vector graphics over bitmap graphics in terms of animation	8
1.1.3. Lack of support on UI/UX creation for complex interactivity	9
1.2. Project Aim	10
1.3. Objectives	10
1.4. Scope	10
1.4.1. Tasks to be executed:	10
1.4.2. Constraints	11
1.4.3. What will be done in this project:	12
1.4.4. What will not be done in this project:	12
1.5. Potential Benefit	12
1.5.1. Tangible Benefit	12
1.5.2. Intangible Benefit	12
1.5.3. Target User	13
1.6. Overview of the IR	13
1.7. Project Plan	13
CHAPTER 2: Literature Review	13
2.1. Domain Research	13
2.2. Global Illumination	13
2.3. Vector Graphics	14
2.4. Interactive UI/UX	14
2.5. Similar Systems/Works	15
2.6. Technical Research	15
CHAPTER 3: Methodology	15
3.1. System Development Methodology	15
3.2. Data Gathering Design	15
3.3. Analysis	15
References	15
Appendices	15

References	16
-------------------------	-----------

Figures

Figure 1: Vector vs Bitmap graphics (Ratermanis, 2017) 8

Tables

CHAPTER 1: INTRODUCTION

Achieving visually rich and interactive content in real-time without compromising performance is a key aspect of immersive gameplay. This project addresses two major challenges in modern game development: creating dynamic, interactive user experiences and implementing accurate, real-time lighting models. Tackling these challenges requires three key innovations: a compute-centric vector graphics renderer, a programmable approach for developing interactive content, and a performant global illumination technique.

Calculating **global illumination** is extremely computationally expensive, as it requires simulating how light bounces off surfaces and interacts with the environment. Ray tracing is an algorithm that calculates these light interactions by tracing lights from the camera into the scene, following their paths as they bounce off surfaces and interact with materials. Each bounce contributes to the final color and lighting of the scene, accounting for reflections, refractions, and scattering.

Unfortunately, ray tracing is too slow for real-time applications, like games. New techniques like light probes and light baking has been employed to approximate global illumination in modern game engines. However, the major issue still exists for these techniques — scalability to larger and more complex scenes.

Vector graphics is a form of computer graphics where visual images are generated from geometric shapes, such as points, lines, curves, and polygons, defined on a Cartesian plane. Vector graphics are often used in situations where scalability and precision are essential. Common applications include: logos, typography, diagrams, charts, motion graphics, etc. Examples of softwares that generates or uses vector graphics content includes Adobe Illustrator, Adobe After Effects, Affinity Publisher, Graphite, and many more. Vector graphics is also used in a wide range of file formats including Scalable Vector Graphics (SVG), Portable Document Format (PDF), TrueType Font (TTF), OpenType Font (OTF), etc. However, these formats are rarely used in the game industry directly (they are often preprocessed into some other formats, i.e. triangulation or signed distance fields [SDF]), as game engines are often only tailored towards rendering triangles and bitmap textures instead of paths and curves that we see in the vector graphics formats.

Markup languages (i.e. Hypertext Markup Language [HTML], Extensible Markup Language [XML]) and style sheets (i.e. Cascading Style Sheets [CSS]) has dominated the way developers layout and style contents. Over the years, technologies like Unity UI Toolkit has evolved in the game industry to adopt the same pattern but with a user friendly editor, allowing users to

layout content using a drag and drop manner while styling their content using sliders, color pickers, and input fields (Jacobsen, 2023). While this improves the user experience of content creation, it lacks the capability of integrating logic and custom contents right inside the user interfaces. These features are often delegated to the programmer which can lead to unintended miscommunications.

1.1. Problem Background

Limitations of scalability in current real-time global illumination techniques

Global illumination has been a notoriously hard problem to solve in computer graphics. To put things into perspective, global illumination intends to solve the *many to many* interactions between light, obstacles, and volumes. In real-time game engines like Unity and Unreal Engine, light probes (a.k.a radiance probes) are placed around the scene to capture lighting information, which can then be applied to nearby objects. To smoothen out the transition between probes, objects interpolate between nearest surrounding probes, weighted by distance to approximate the global radiance.

This technique leads to questions like “how many probes should a scene have?” or “how much probes is a good approximation?”. Ultimately, it becomes a trade-off between fidelity versus performance, with more probes resulting in better approximation, while fewer probes improve performance. This paradoxical issue raises the challenge of finding the optimal balance. This dilemma underscores the need for smarter, adaptive techniques, ensuring both visual fidelity and efficiency.

Advantages of vector graphics over bitmap graphics in terms of animation

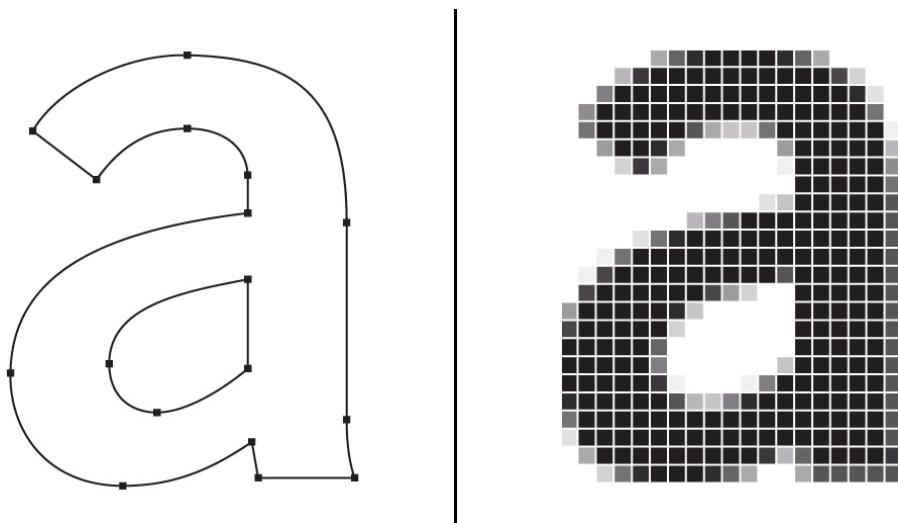


Figure 1: Vector vs Bitmap graphics (Ratermanis, 2017)

Traditional methods of rendering 2D graphics has always relied on bitmap-based texture mapping (Ray et al., 2005). While this approach is ubiquitous, it suffers a major drawback of the *pixelation* effect when being scaled beyond the original resolution (Nehab & Hoppe, 2008). Furthermore, creating animations using bitmap graphics can be extremely limited and complex because of the rigid grid-like data structure used to store the data. Animating bitmap graphics are commonly done through the use of shaders which directly manipulates the individual pixels, or relying on image sequences (flipbooks) which produces an illusion of movement.

Unlike raster graphics, which rely on a fixed grid of pixels, vector graphics are resolution-independent. This means that it can scale without losing quality (shown in Figure 1). A vector illustration is composed of multiple *paths* that define *shapes* to be painted in a given order (Ganacim et al., 2014). Each of these individual paths can be traced, altered, or even morphed into a completely different shape which allows for a huge variety of animation techniques.

Lastly, it is crucial to recognize that while vector graphics offer numerous benefits, it is only suitable for representing precise shapes — such as fonts, logos, and icons. In contrast, complex images with intricate details, like photographs of a cat are far better represented using bitmap formats.

Lack of support on UI/UX creation for complex interactivity

Most game engines in the market like Unity, Godot, Game Maker, and Unreal Engine uses a *What You See Is What You Get* (WYSIWYG) editor for creating user interfaces. WYSIWYG editors are visual centric tools that let users work directly within the presentation form of the content (Mädje, 2022). Users normally layout their contents using a drag and drop editor and then style them using a style-sheet. To bind interactions or animations towards a content, users would need to label it with a unique tag and query them through code.

Complex content and logic wouldn't be possible through a typical WYSIWYG editor. For instance, it is virtually impossible to author a custom polygon shape in the editor with custom math based animation which depends on a time value. This can only be achieved through code, and is often limited by the application programming interface (API) layer provided by the WYSIWYG editor. This creates a huge distinction between the game/UI logic and the visual representation that is needed to convey the messages.

While hot-reloading is applicable for the layout and styling (and simple logic to some extend) of a content. A WYSIWYG editor would not be capable of hot-reloading complex logic as these

can only be achieved using code, which in most cases, requires a re-compilation. This could lead to frustration and lost of creativity due to the slow feedback loop.

1.2. Project Aim

This project aims to empower creators to create rich and visually appealing content in games in an efficient and streamlined workflow, by allowing them to focus most of their time on the content instead of the technical details needed to achieve the look or feel that they envisioned.

1.3. Objectives

The objectives of this project are as follows:

1. To utilize Vello, a compute-centric vector graphics renderer for rendering animated and dynamic vector graphics content.
2. To create an intuitive and yet powerful (programmable) workflow for generating animated and dynamic content using Typst.
3. To allow creators to focus on the creative aspects of game development.
4. To implement radiance cascades, a technique that provides realistic lighting without sacrificing real-time performance.

1.4. Scope

The scope of this project involves making a game that utilizes 2 custom libraries (1 for global illumination and 1 for interactive vector graphics content). The creation of the game will ensure that 2 of our libraries are production ready by the end of this project.

Libraries (Crates)	
Bevy Radiance Cascades (Bevy RC)	A 2D global illumination solution for the Bevy game engine.
Velyst	An interactive Typst content creator using Vello and Bevy.
Game	
Lumina	A 2D top down fast paced objective based PvPvE game.

Tasks to be executed:

1. Develop the **Bevy RC** crate.
 - a. Develop the radiance cascades algorithm.
 - b. Implement radiance cascades into Bevy's 2D render graph.
 - c. Support emissive and translucent materials.
 - d. Support directional light beams / spot lights.

- e. Support negative lighting effects (light consumption).
 - f. Support rim lighting for better visual effects.
2. Develop the **Velyst** crate.
- a. Develop an integrated compiler for Typst content in Bevy.
 - b. Support hot-reloading of Typst content.
 - c. Support interactivity between Bevy and Typst.
 - d. Develop an easy-to-use workflow for UI creation using Typst.
 - e. Write up a getting started documentation to make on-boarding easier for new developers.
3. Develop the **Lumina** game.
- a. Create a game design document (GDD) for the game.
 - b. Integrate both **Bevy RC** and **Velyst** into the game.
 - c. Develop all the required game mechanics for the game.
 - d. Playtest the game and gather player feedbacks on the game.

Constraints

Constraint	Reason
Compatibility	Our project uses multiple cutting-edge and innovative technologies. This means that some of the technologies we depend on might be experimental or have yet to stabilize. This makes it difficult to ensure cross platform / device compatibility for the systems we built. For example, Vello requires compute shaders to render vector graphics, which means it can only run on newer versions of browsers that support WebGPU.
Limited Documentation	Because some of the technologies we use are experimental or less widely adopted, available documentation and community support may be limited. This slows down the development process as issues can be difficult to troubleshoot without established resources.
Limited Resource	The project is subject to limitations in terms of budget, personnel, and time. Allocating sufficient resources to develop, test, and refine the workflow for real-time global illumination and vector graphics is crucial. Any constraints in these areas can impact the project's scope and delivery timeline.

What will be done in this project:

1. **Global illumination:** Our implementation of radiance cascades will support any kind of 2D shapes with emissive, translucent, and non-emissive color material. It is aimed to be a plug and play plugin. Users should be able to use it without the need to learn about complex configurations.
2. **Typst compiler:** A custom implementation of Typst compiler will be created to fit the purpose of real-time Typst content rendering. This compiler should be able to re-compile Typst content on demand, allowing developers to view reflect their saved changes immediately.
3. **Dynamic vector graphics:** The **Velyst** crate will allow users to generate dynamic vector graphics content which is then rendered through the Vello renderer.
4. **Game demo prototype:** **Lumina** will be developed to showcase all of the above in a compact game format. Players will experience realistic and beautiful 2D lighting, as well as interactive vector graphics UI elements.

What will not be done in this project:

1. **No reflections & refractions:** Radiance cascades by default does not support reflections and refractions. Implementing these 2 features requires a tremendous amount of additional research, experimentation, and effort. Not to mention that implementing radiance cascades itself can already be extremely challenging as an individual working in such time constraint.
2. **Not creating an animation library:** An animation library involves preparing a huge variety of common animation effects. This takes a huge amount of time which does not fit the goal and scope of this project. To understand more about animation libraries, we strongly encourage you to look into the *Bevy MotionGfx* project.
3. **Not a commercial game:** Our goal with **Lumina** is to create a game that demonstrates our technologies — **Velyst** and **Bevy RC**. It is not meant to be on par with a full on commercial game.

1.5. Potential Benefit

Tangible Benefit

Intangible Benefit

Target User

Velyst will particularly be targeted towards UI/UX developers, motion graphics creators, and vector graphics enthusiasts. **Bevy RC** will be targeted towards visual effects (VFX) artists, game programmers, and technical artists. Because these technologies are built on top of Rust and Bevy, the general users will come from the Rust + Bevy community.

As for **Lumina**, the target audience are gamers who love fast paced multipalyer games like *Apex Legends* and *Astro Duel 2*. It will particularly appeal to gamers who love the mix of competitive PvP and PvE like *Destiny 2*'s Gambit game mode and *World War Z*.

1.6. Overview of the IR

1.7. Project Plan

CHAPTER 2: Literature Review

2.1. Domain Research

2.2. Global Illumination

As mentioned previously, ray tracing is the de-facto standard for calculating light bounces which contributes to global illumination. Clever methods like backwards ray tracing has been introduced to speed up the algorithm, but it is still nowhere near real-time frame rates (Arvo & others, 1986). Light baking is introduced to solve this issue, however, it lacks the ability to adapt to runtime scene changes.

Recent studies has shown great results of utilizing neural networks for approximating global illumination (Choi et al., 2024). However, neural network based methods tend to suffer from unpredictability as the output is highly basd upon the input training data, making it unreliable.

Recent works by McGuire et al. (2017) places light field probes around the scene to encode lighting information from static objects and sample them in real-time. Dynamic diffuse global illumination (DDGI) further improves this technique by allowing light field probes to update dynamically based on scene changes (Majercik et al., 2019).

Radiance cascades improves upon this technique by using a hierarchical structure to place light probes (Osborne & Sannikov, 2024). This technique is based upon the *penumbra condition*, where closer distance require low angular resolution and high spatial resolution while further distance require high angular resolution and low spatial resolution.

2.3. Vector Graphics

Scanline rendering is the process of shooting rays from one side of the screen to the other while coloring pixels in between based on collision checkings with paths in between. A GPU based scanline rasterization method is proposed by parallelizing over *boundary fragments* while bulk processing non-boundary fragments as horizontal spans (Li et al., 2016). This method allows fully animated vector graphics to be rendered in interactive frame rates.

Apart from scanline rasterization, tessellation method can also be used to convert vector graphics into triangles and then pushed to the GPU for hardware accelerated rasterization. Loop & Blinn (2005) further improved this method by removing the need of approximating curve segments into lines. Instead, each curve segments is evaluated in a *fragment shader* which can be calculated on the GPU. This allows for extreme zoom levels without sacrificing qualities.

Re-tesselation of vector graphics can be computationally expensive, especially when it's inherently a serial algorithm that often needs to be solved on the CPU. Kokojima et al. (2006) combines the work of Loop & Blinn (2005) with the usage of GPU's stencil buffer by using *triangle fans* to skip the tessellation process. This approach, however, does not extend to cubic Bézier segments as they might not be convex. Rueda et al. (2008) addressed this issue by implementing a fragment shader that evaluates the implicit equation of the Bézier curve to discard the pixels that fall outside it. The two-step "Stencil then Cover" (StC) method builds upon all of these work and unified path rendering with OpenGL's shading pipeline — `NV_path_rendering` (Kilgard & Bolz, 2012). This library was further improved upon by adding support for transparency groups, patterns, gradients, more color spaces, etc. (Batra et al., 2015). It was eventually integrated into Adobe Illustrator.

2.4. Interactive UI/UX

Beneath all graphical interfaces lies the underlying code that structures and renders the visual elements. The two most notable approaches towards creating user interface frameworks are immediate-mode graphical user interface (IMGUI) and retained-mode graphical user interface (RMGUI). Some popular IMGUI frameworks include Dear IMGUI and egui (*Dear ImGui*, n.d.; *Egui*, n.d.), while some popular RMGUI frameworks include Xilem (*Xilem*, n.d.). Although powerful, these UI frameworks strongly rely on hardcoded programming.

Enter the web technologies. Modern browsers typically render UI elements using markup languages like HTML and SVG for structuring the content and style-sheets like CSS for styling them. The use of markup structures allows developers to fully separate their UI layout from the codebase, simplifying the identification and management of UI components. With style sheets,

developers can create, share, and reuse templates, enhancing consistency and streamlining the design process throughout the application. Notable frameworks that utilize this model include Unity UI Toolkit, React, Vue, etc (Jacobsen, 2023; *React*, n.d.; *Vue*, n.d.).

Markup languages also give rise to many WYSIWYG editors. These editors let users perform drag and drop actions to layout UI for quick prototyping as each component can now be represented using only markup syntax (no code required).

A major limitation of simple markup languages like HTML is that structure changes can only be achieved through code. For example, if you want a form to disappear after button press, you would need to alter the HTML via code. Typst offers an alternative towards this problem by introducing programming capabilities into markdown.

Typst is a competitor of LaTeX, designed to simplify the typesetting process with a modern and intuitive approach. Unlike its predecessors, Typst can directly embed logic. Using the previous example, developers would only need to pass in a boolean value and Typst will automatically exclude the form from being in the layout at all. This currently works only in theory, as Typst is primarily a document generator without a user-friendly interface for modifying defined variables.

This is where our project comes in, we aim to provide this interface through Velyst, which couples Typst with Vello for rendering dynamic and programmable content in real-time.

2.5. Similar Systems/Works

2.6. Technical Research

CHAPTER 3: Methodology

3.1. System Development Methodology

3.2. Data Gathering Design

3.3. Analysis

References

Appendices

References

- Arvo, J., & others. (1986). Backward ray tracing. *Developments in Ray Tracing, Computer Graphics, Proc. Of ACM SIGGRAPH 86 Course Notes*, 259–263.
- Batra, V., Kilgard, M. J., Kumar, H., & Lorach, T. (2015). Accelerating vector graphics rendering using the graphics hardware pipeline. *ACM Transactions on Graphics (TOG)*, 34(4), 1–15.
- Choi, E., Carpentier, V., Shin, S., & Yoo, S. (2024). Baking Relightable NeRF for Real-time Direct/Indirect Illumination Rendering. *Arxiv Preprint Arxiv:2409.10327*.
- Dear ImGui*. Retrieved September 22, 2024, from <https://github.com/ocornut/imgui>
- Egui*. Retrieved September 22, 2024, from <https://github.com/emilk/egui>
- Ganacim, F., Lima, R. S., De Figueiredo, L. H., & Nehab, D. (2014). Massively-parallel vector graphics. *ACM Transactions on Graphics (TOG)*, 33(6), 1–14.
- Jacobsen, T. K. (2023, November). *New UI toolkit demos for programmers and artists: Unity blog*. Unity. <https://unity.com/blog/engine-platform/new-ui-toolkit-demos-for-programmers-artists>
- Kilgard, M. J., & Bolz, J. (2012). GPU-accelerated path rendering. *ACM Transactions on Graphics (TOG)*, 31(6), 1–10.
- Kokojima, Y., Sugita, K., Saito, T., & Takemoto, T. (2006). Resolution independent rendering of deformable vector objects using graphics hardware. In *ACM SIGGRAPH 2006 Sketches: ACM SIGGRAPH 2006 Sketches* (p. 118–es).
- Li, R., Hou, Q., & Zhou, K. (2016). Efficient GPU path rendering using scanline rasterization. *ACM Transactions on Graphics (TOG)*, 35(6), 1–12.
- Loop, C., & Blinn, J. (2005). Resolution independent curve rendering using programmable graphics hardware. In *ACM SIGGRAPH 2005 Papers: ACM SIGGRAPH 2005 Papers* (pp. 1000–1009).
- Majercik, Z., Guertin, J.-P., Nowrouzezahrai, D., & McGuire, M. (2019). Dynamic diffuse global illumination with ray-traced irradiance fields. *Journal of Computer Graphics Techniques*, 8(2).

McGuire, M., Mara, M., Nowrouzezahrai, D., & Luebke, D. (2017). Real-time global illumination using precomputed light field probes. *Proceedings of the 21st ACM SIGGRAPH Symposium on Interactive 3d Graphics and Games*. <https://doi.org/10.1145/3023368.3023378>

Mädje, L. (2022). A Programmable Markup Language for Typesetting. *Technical University of Berlin*, 1–77.

Nehab, D., & Hoppe, H. (2008). Random-access rendering of general vector graphics. *ACM Transactions on Graphics (TOG)*, 27(5), 1–10.

Osborne, C. M., & Sannikov, A. (2024). Radiance Cascades: A Novel High-Resolution Formal Solution for Multidimensional Non-LTE Radiative Transfer. *Arxiv Preprint Arxiv:2408.14425*, 1–21.

Ratermanis, A. (2017, December 29). *Vector vs raster: What's best for your logo*. <https://www.ratermanis.com/blog/2017/12/28/vector-vs-raster>

Ray, N., Cavin, X., & Lévy, B. (2005). Vector Texture Maps on the GPU. *Inst. ALICE (Algorithms, Comput., Geometry Image Dept. INRIA Nancy Grand-Est/loria), Tech. Rep. ALICE-TR-05-003*.

React. Retrieved September 22, 2024, from <https://react.dev/>

Rueda, A., De Miras, J. R., & Feito, F. R. (2008). GPU-based rendering of curved polygons using simplicial coverings. *Computers & Graphics*, 32(5), 581–588.

Vue. Retrieved September 22, 2024, from <https://vuejs.org/>

Xilem. Retrieved September 22, 2024, from <https://github.com/linebender/xilem>